



Enarx

Protection for data in use

Mike Bursell
Office of the CTO

Nathaniel McCallum
Sr. Principal Engineer


<https://enarx.io>

The Problem

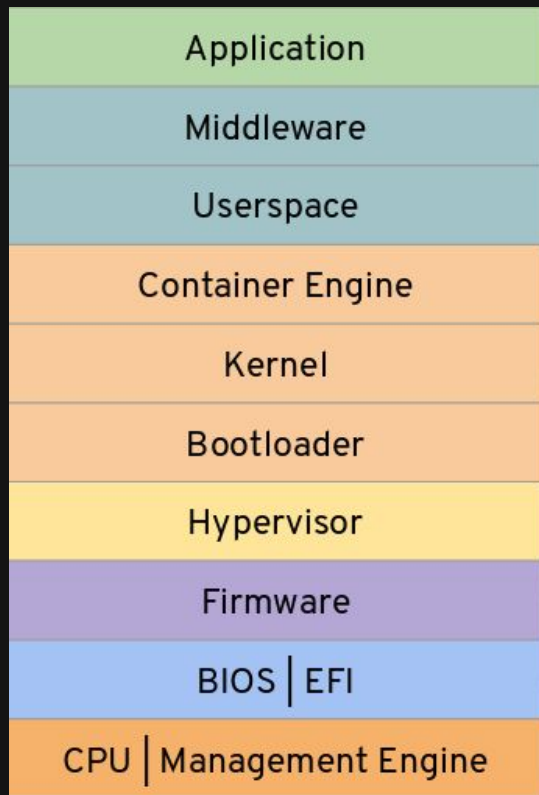
The Need for Confidentiality and Integrity

- Banking & Finance
- Government & Public Sector
- Telco
- IoT
- HIPAA
- GDPR
- Sensitive enterprise functions
- Defense
- Human Rights NGOs
- ...

Virtualization Stack




Container Stack



<https://xkcd.com/2166/>

THE MODERN TECH STACK



The Plan



The Principles

Don't trust the **host**
Don't trust the host **owner**
Don't trust the host **operator**
All **hardware** cryptographically
verified
All **software** audited and
cryptographically verified



The Fit

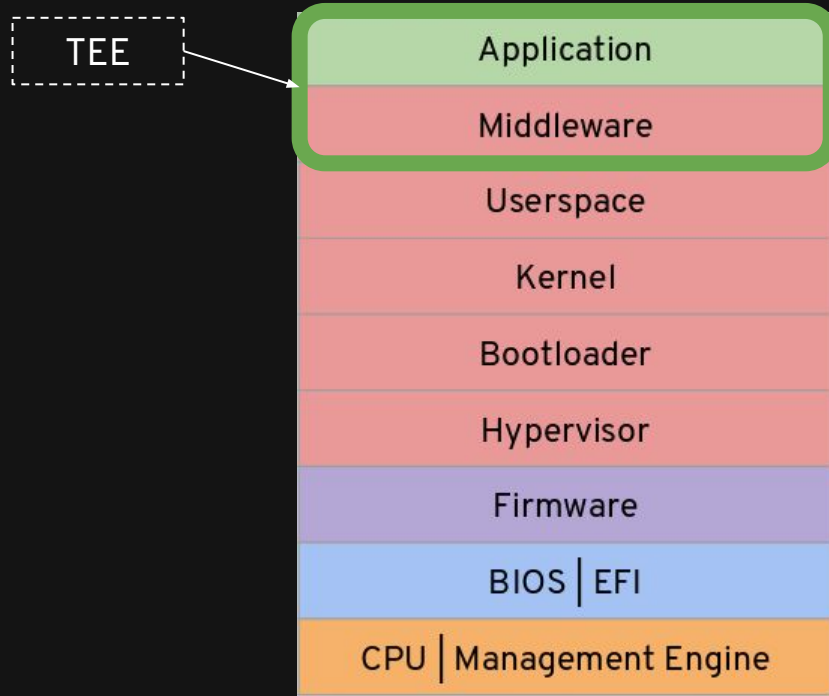
Don't trust the **host**
Don't trust the host **owner**
Don't trust the host **operator**
All **hardware** cryptographically
verified
All **software** audited and
cryptographically verified



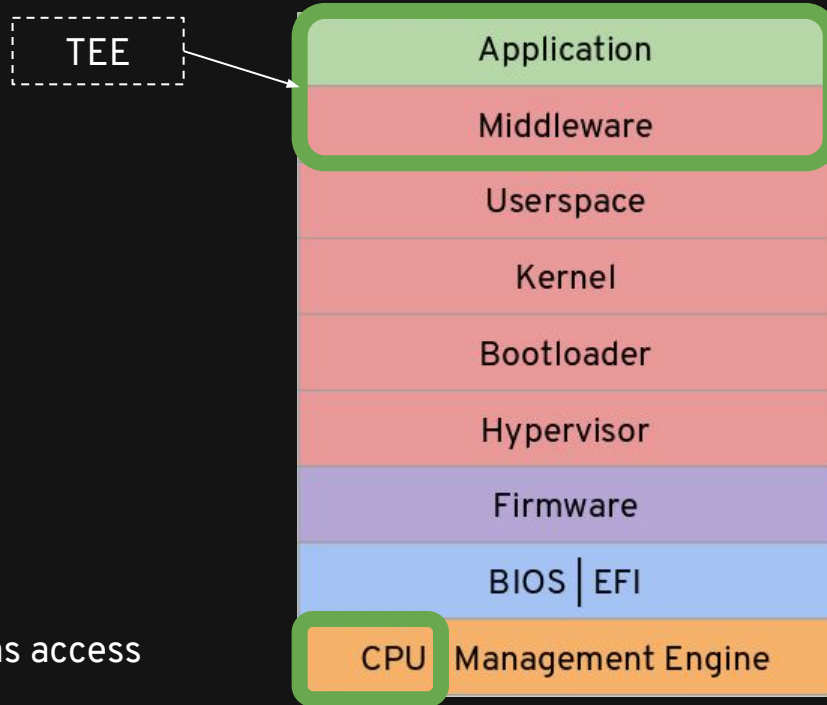
Well suited to **microservices**
Well suited to **sensitive data or**
algorithms
Easy **development integration**
Simple **deployment**
Standards based: **WebAssembly**
(WASM)

Trusted Execution Environments

What's a TEE?



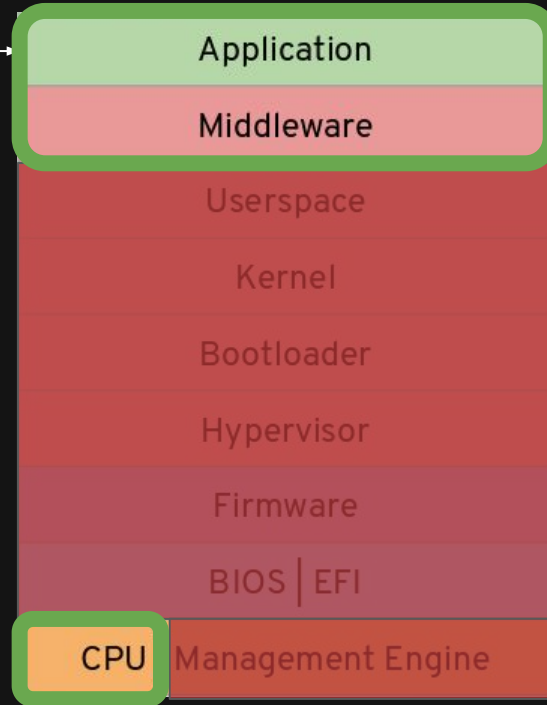
What's a TEE?



Only the CPU has access

What's a TEE?

TEE

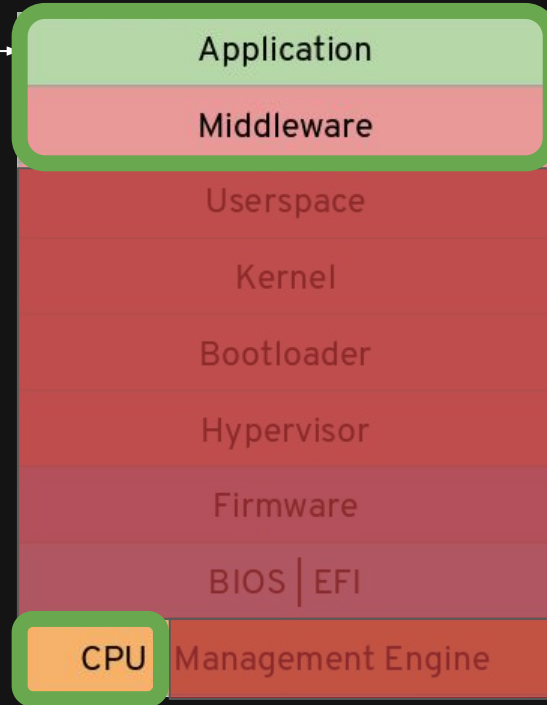


What happens when
other layers try to
access?

Only the CPU has access

What's a TEE?

TEE



What happens when
other layers try to
access?

Blocked by CPU.

Only the CPU has access

Trusted Execution Environments



TEE is a protected area within the host, for execution of sensitive workloads

Trusted Execution Environments



TEE is a protected area within the host, for execution of sensitive workloads


TEE provides:

- Memory Confidentiality
- Integrity Protection
- General compute
- HWRNG

Trusted Execution Environments



Q. “But how do I know that it’s a valid TEE?”



TEE provides:

- Memory Confidentiality
- Integrity Protection
- General compute
- HWRNG

Trusted Execution Summary



Q. “But how do I know that it’s a valid TEE?”

A. Attestation

TEE provides:

- Memory Confidentiality
- Integrity Protection
- General compute
- HWRNG

Trusted Execution Summary



Attestation includes:

- Diffie-Hellman Public Key
- Hardware Root of Trust
- TEE Measurement

TEE provides:

- Memory Confidentiality
- Integrity Protection
- General compute
- HWRNG

Trusted Execution Summary



Attestation includes:

- Diffie-Hellman Public Key
- Hardware Root of Trust
- TEE Measurement

TEE provides:

- Memory Confidentiality
- Integrity Protection
- General compute
- HWRNG

Trusted Execution Models

Process-Based

- Intel SGX (not upstream)
- RISC-V Sanctum (no hardware)

VM-Based

- AMD SEV
- IBM PEF (no hardware)
- Intel MKTME (no attestation¹)

Not a TEE: TrustZone, TPM

1. Attestation is discussed here: <https://patents.google.com/patent/US20190042463A1/en?q=20190042463>

Trusted Execution: Process-Based

PROS

- Access to system APIs from Keep

CONS

- Unfiltered system API calls from Keep
- Application redesign required
- Untested security boundary
- Fantastic for malware
- Lock-in

Trusted Execution: Virtual Machine-Based

PROS

- Strengthening of existing boundary
- Run application on existing stacks
- Bidirectional isolation
- Limits malware

CONS

- Hardware emulation
- Heavy weight for microservices
- CPU architecture lock-in
- Duplicated kernel pages
- Host-provided BIOS

Open hybrid cloud and Enarx

Step 1: on premises

Internal

Internet

Trusted

Semi-trusted

Untrusted

—

- -

. . . .

Internal dev

Step 1: on premises

Internal

Internet

Trusted

Semi-trusted

Untrusted

—


- -

. . . .


Owned host

Internal dev


Step 2: private cloud




Step 2: private cloud



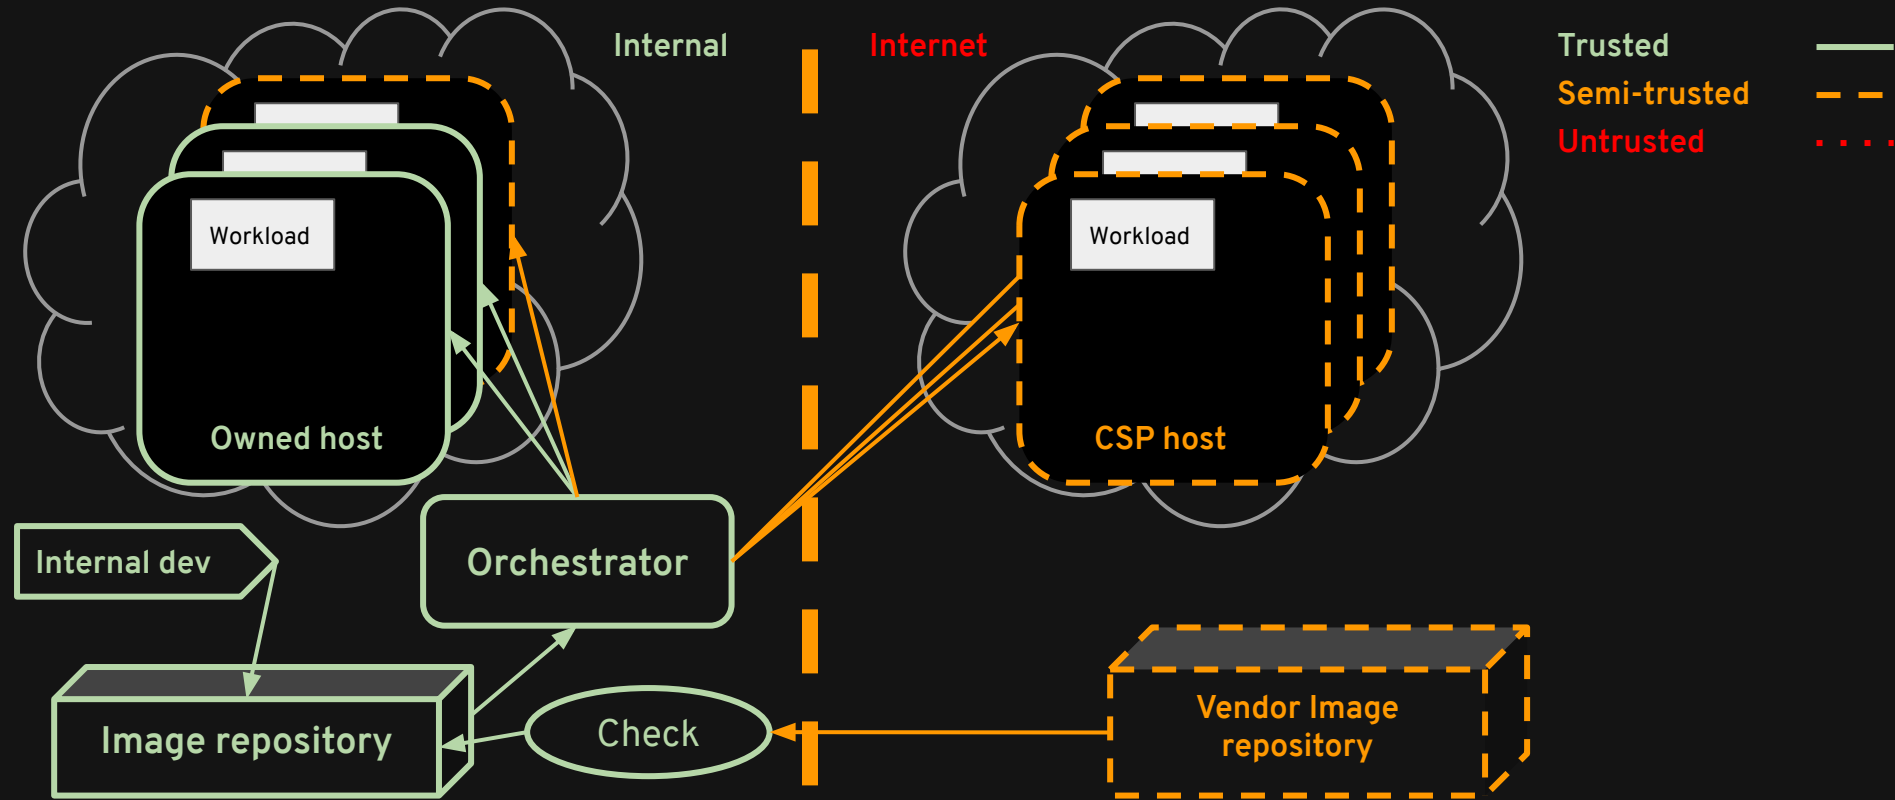
Step 2: private cloud



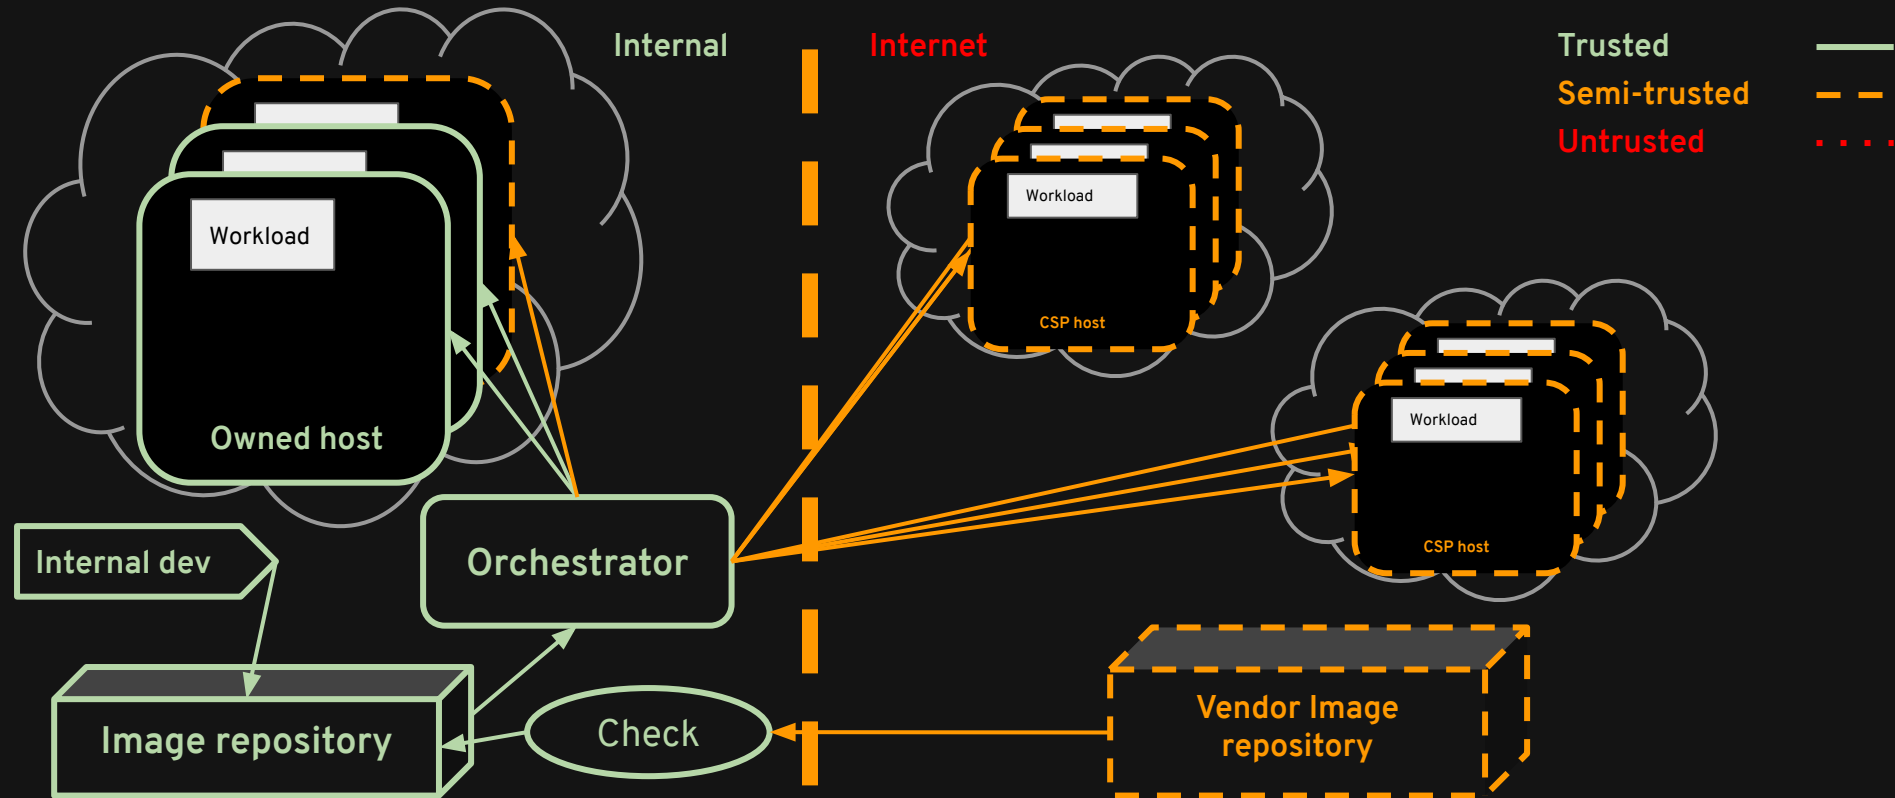
Step 3: public cloud



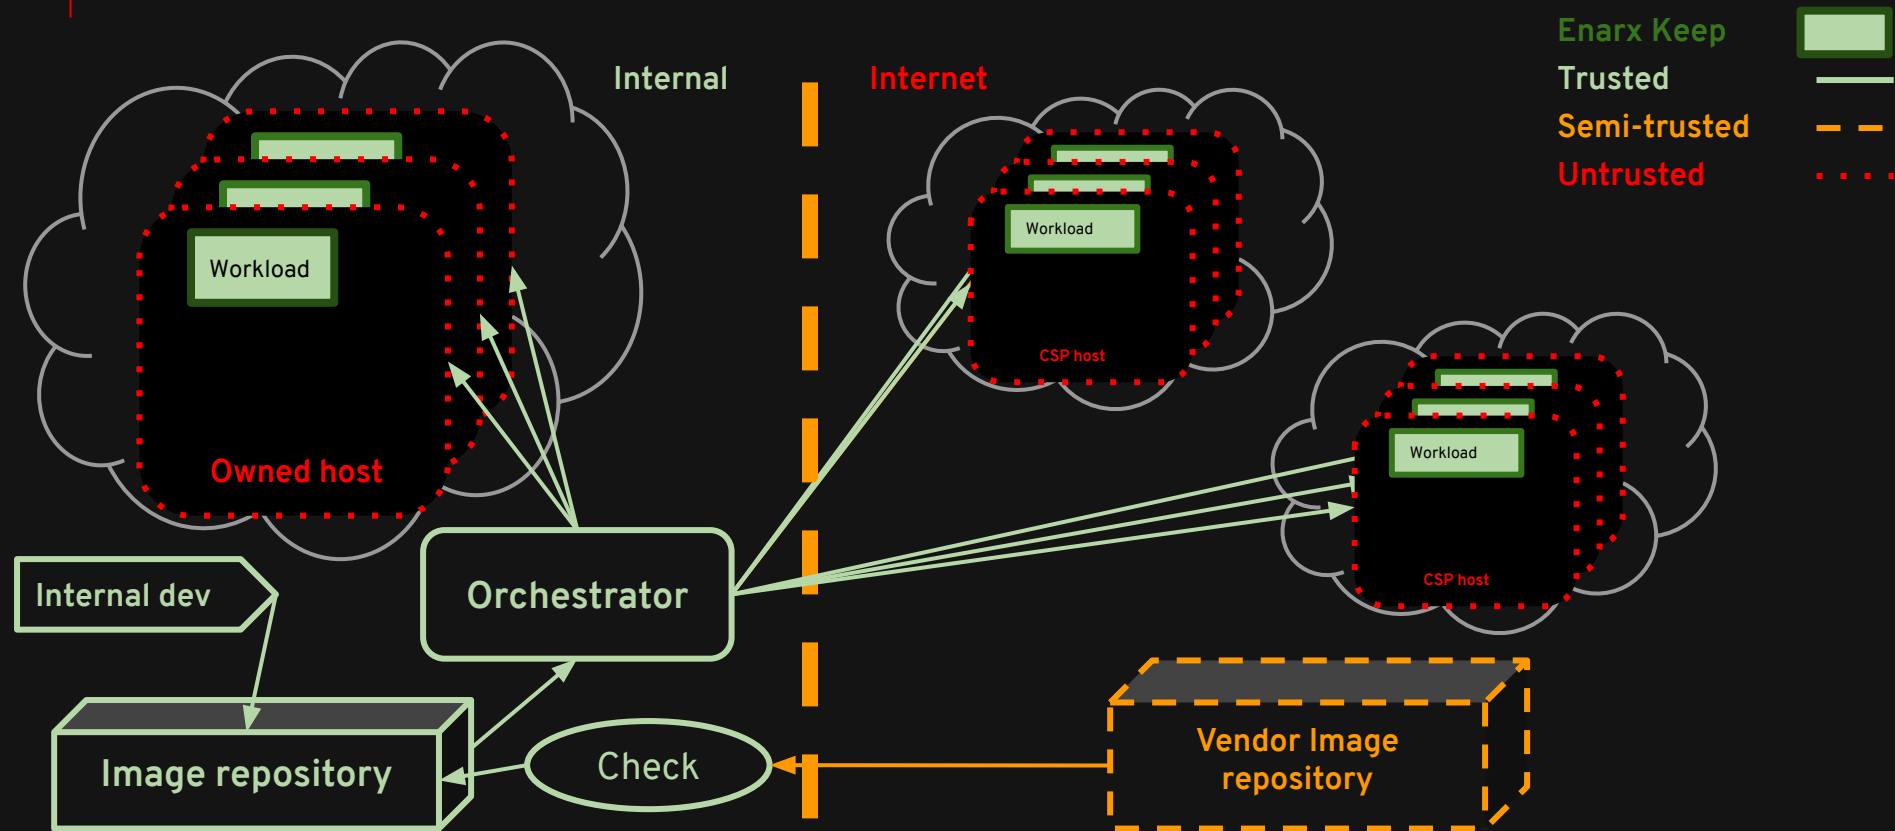
Step 4: hybrid cloud



Step 5: hybrid multicloud

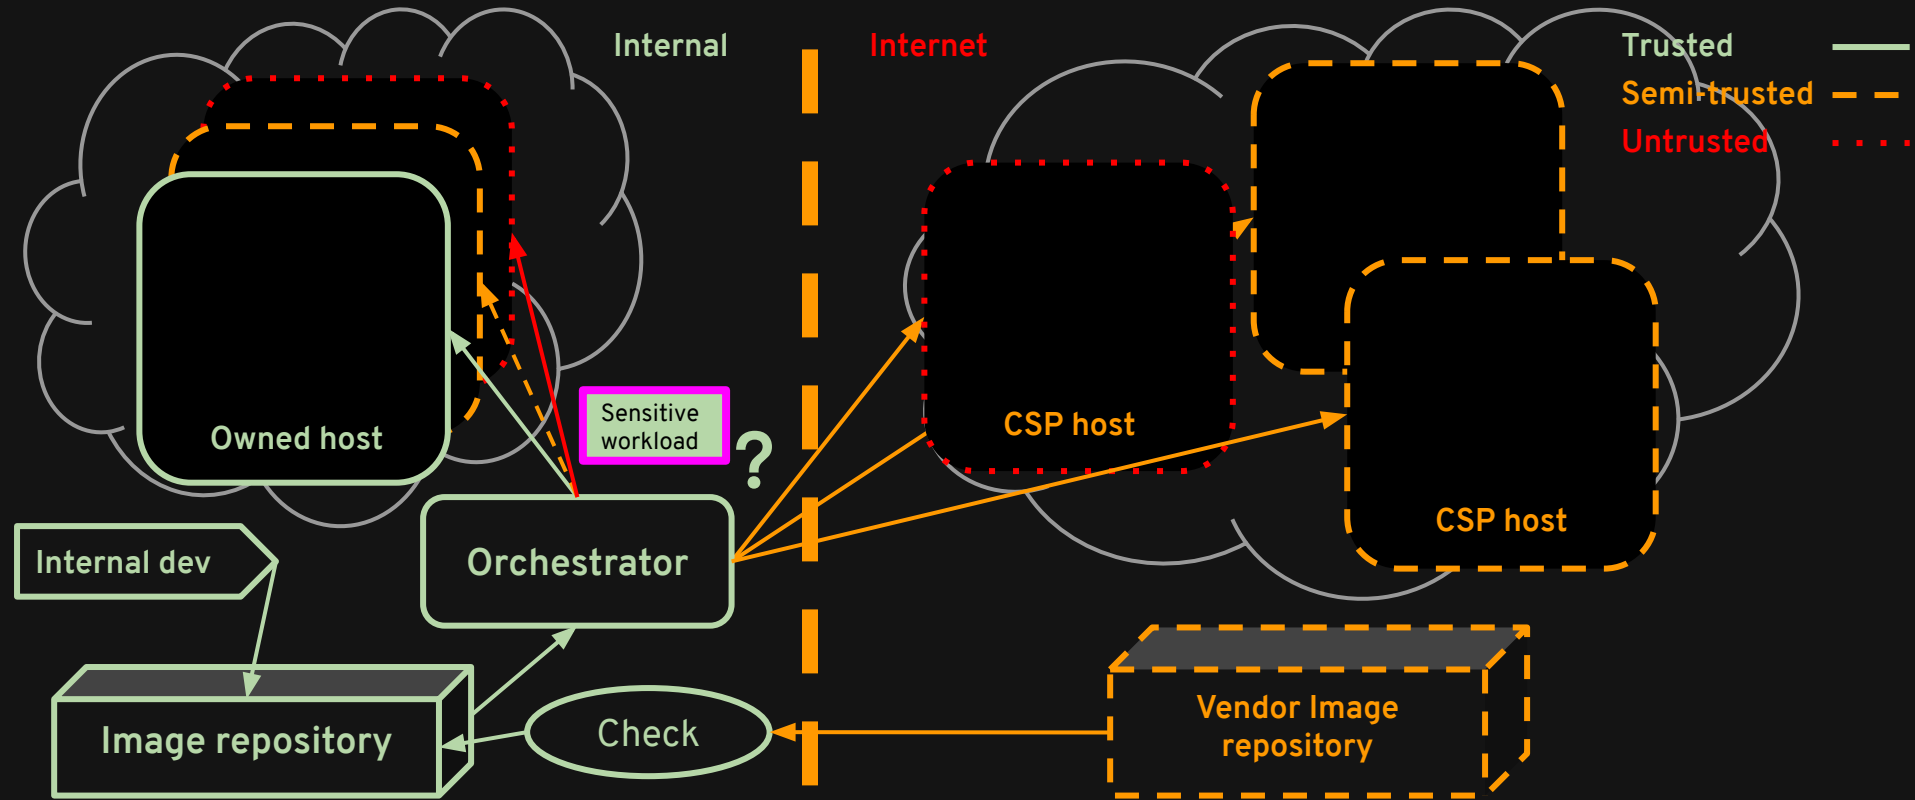


Step 6: Enarx hybrid multicloud




New options for workloads with Enarx


Mix and match for different workload types & Enarx




Mix and match for different workload types & Enarx




Mix and match for different workload types & Enarx



Mix and match for different workload types & Enarx



Mix and match for different workload types & Enarx



New options for orchestration with Enarx

Mix and match for different workload types & Enarx

POSSIBLE MODEL

Internal

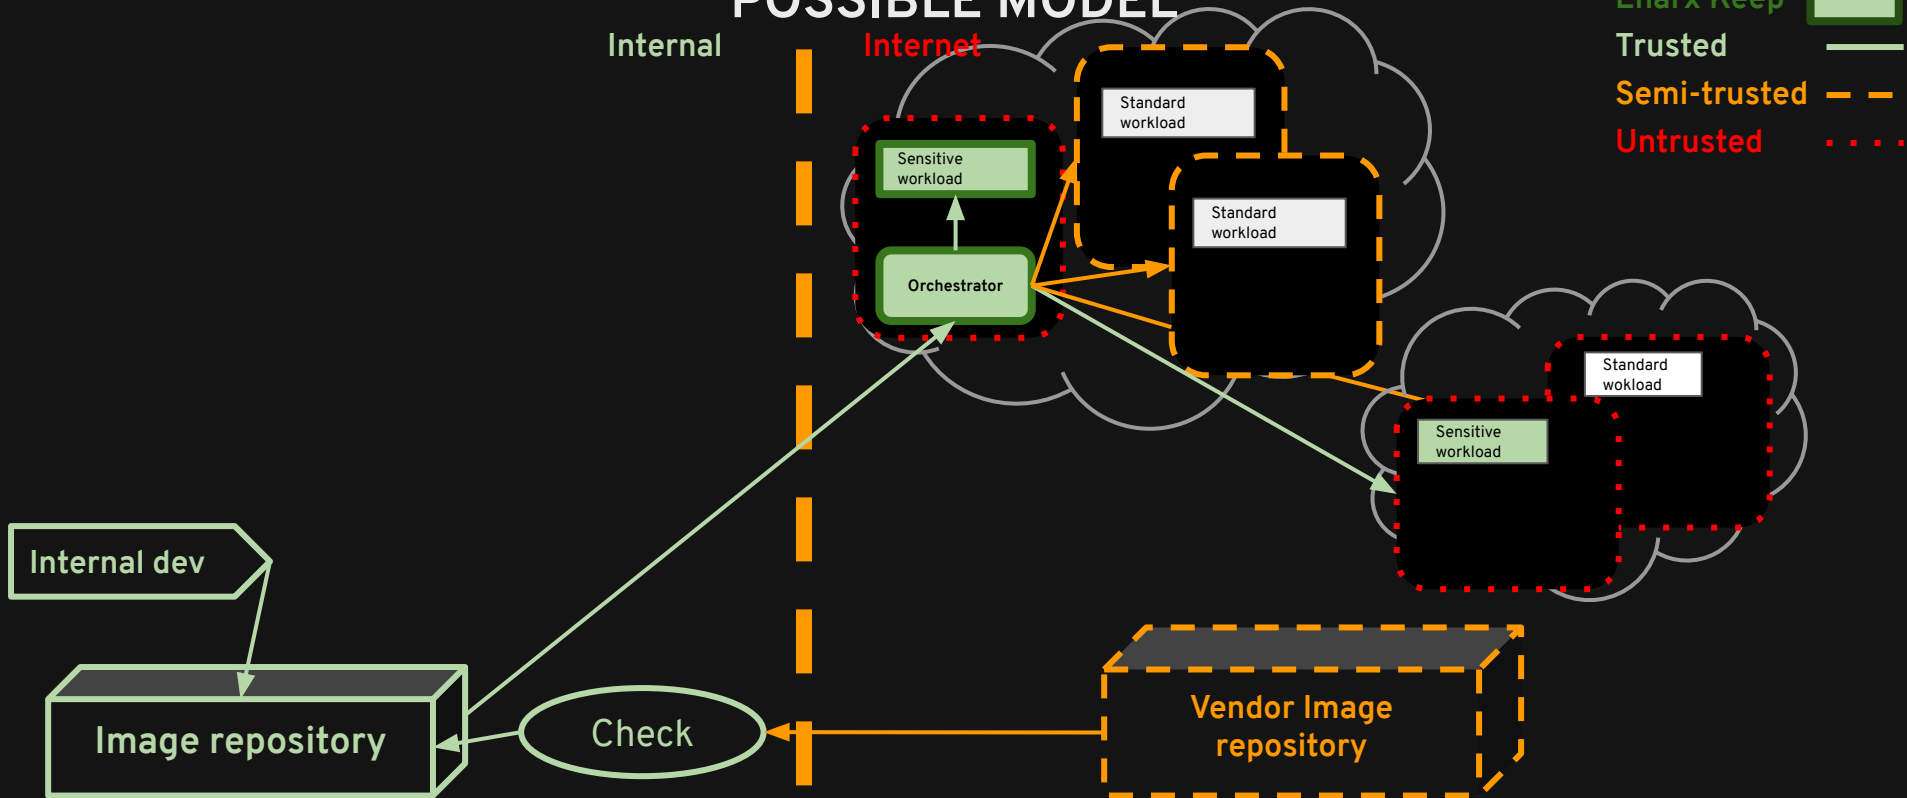
Internet

Enarx Keep

Trusted

Semi-trusted

Untrusted



Value of the trusted stack

How trusted stacks help

Hardware platform

How trusted stacks help

RHEL

Hardware platform


How trusted stacks help

Openshift


RHEL

Hardware platform

How trusted stacks help

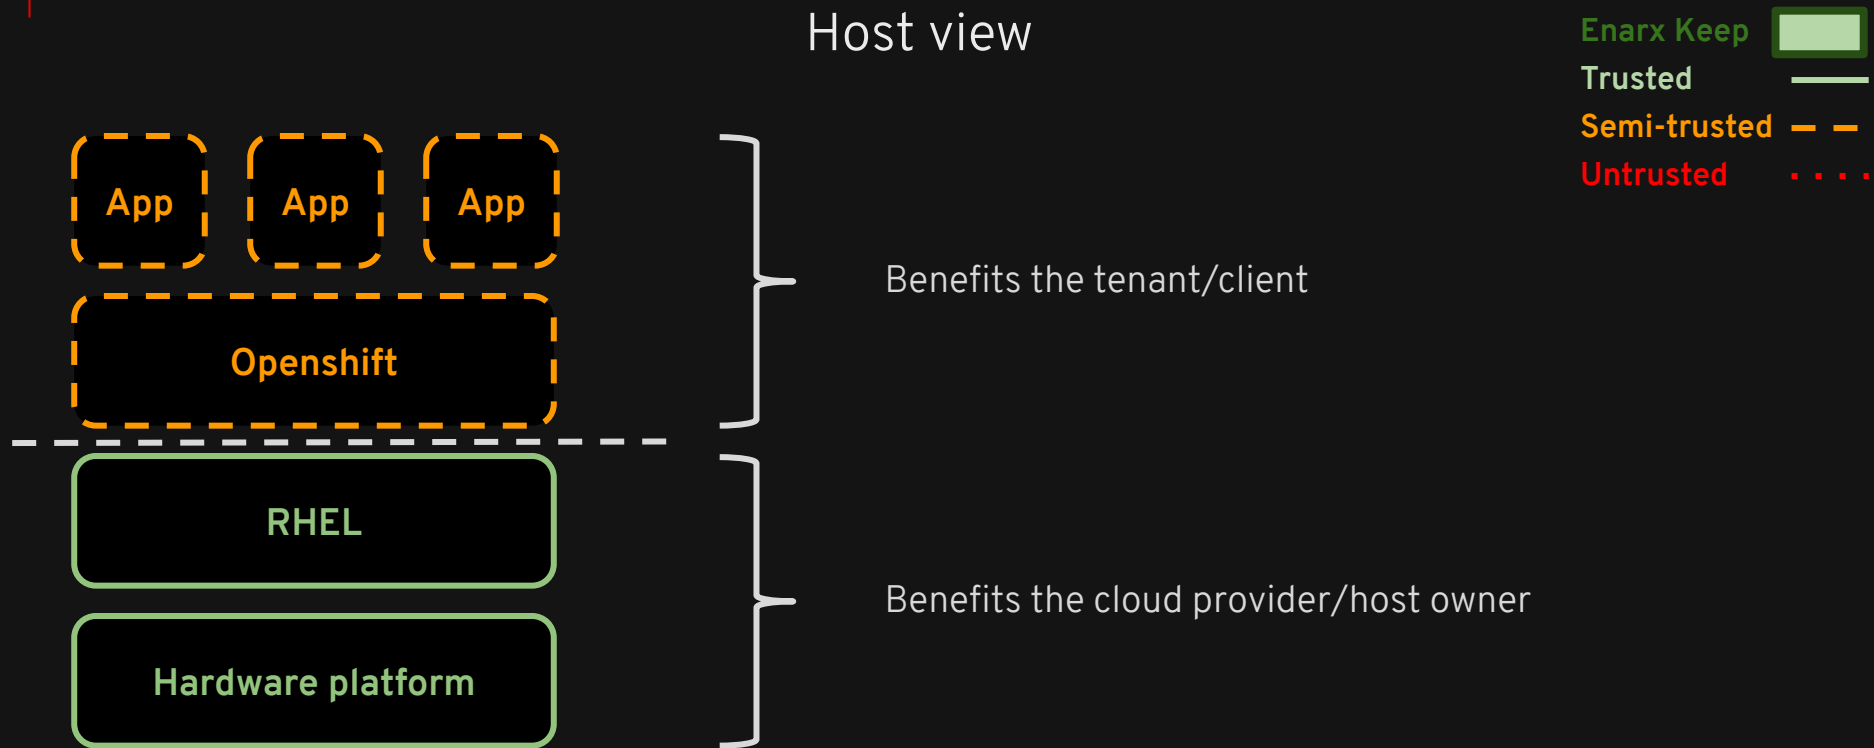


How trusted stacks help




How trusted stacks help

Host view



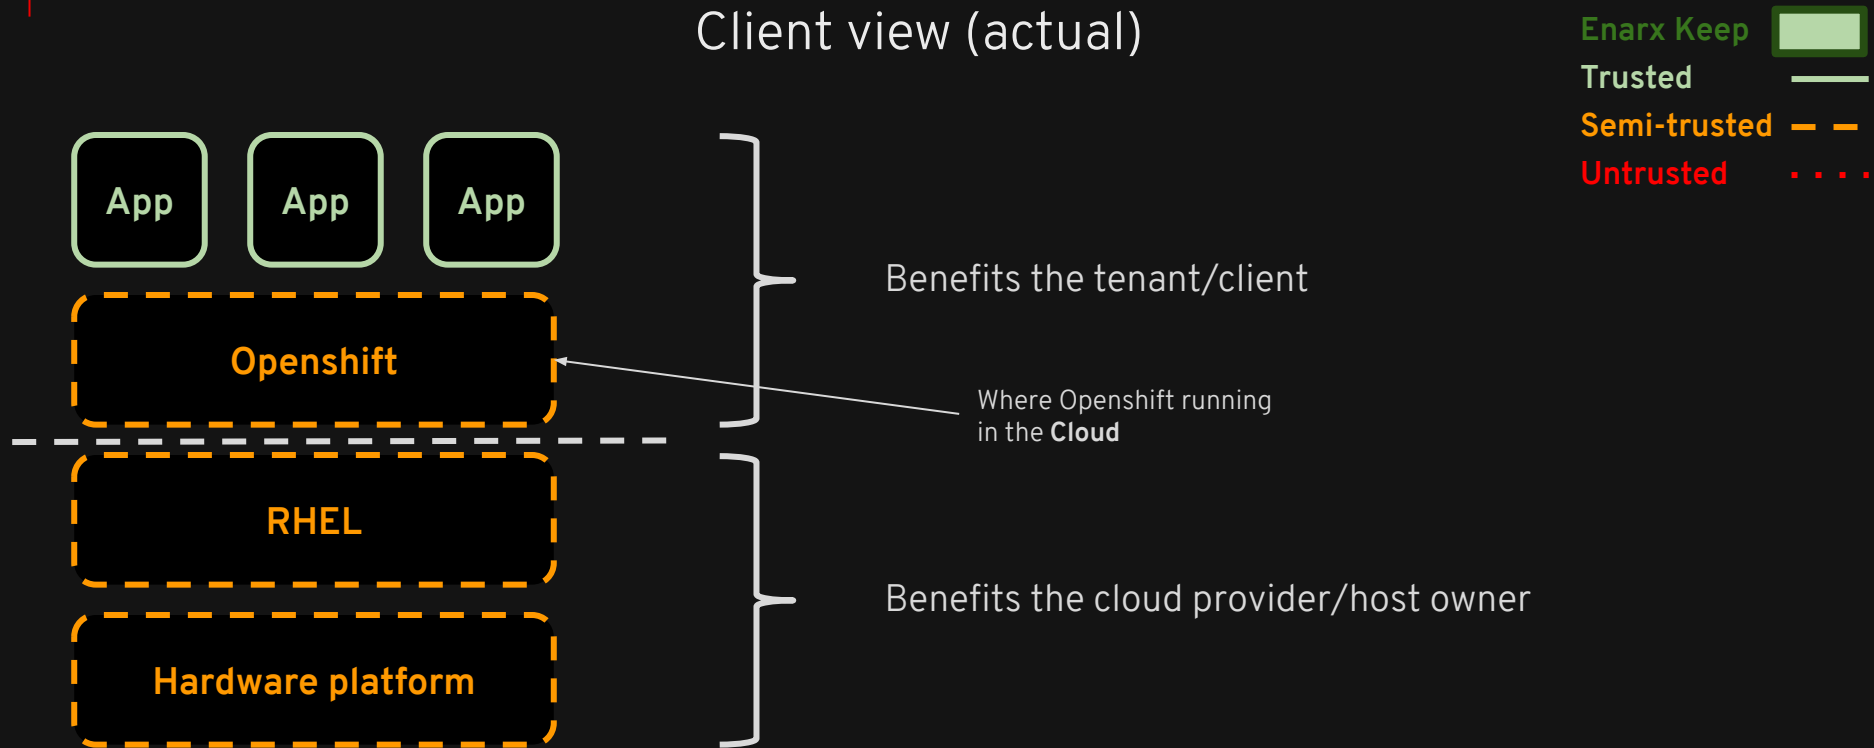
How trusted stacks help

Client view (hoped)



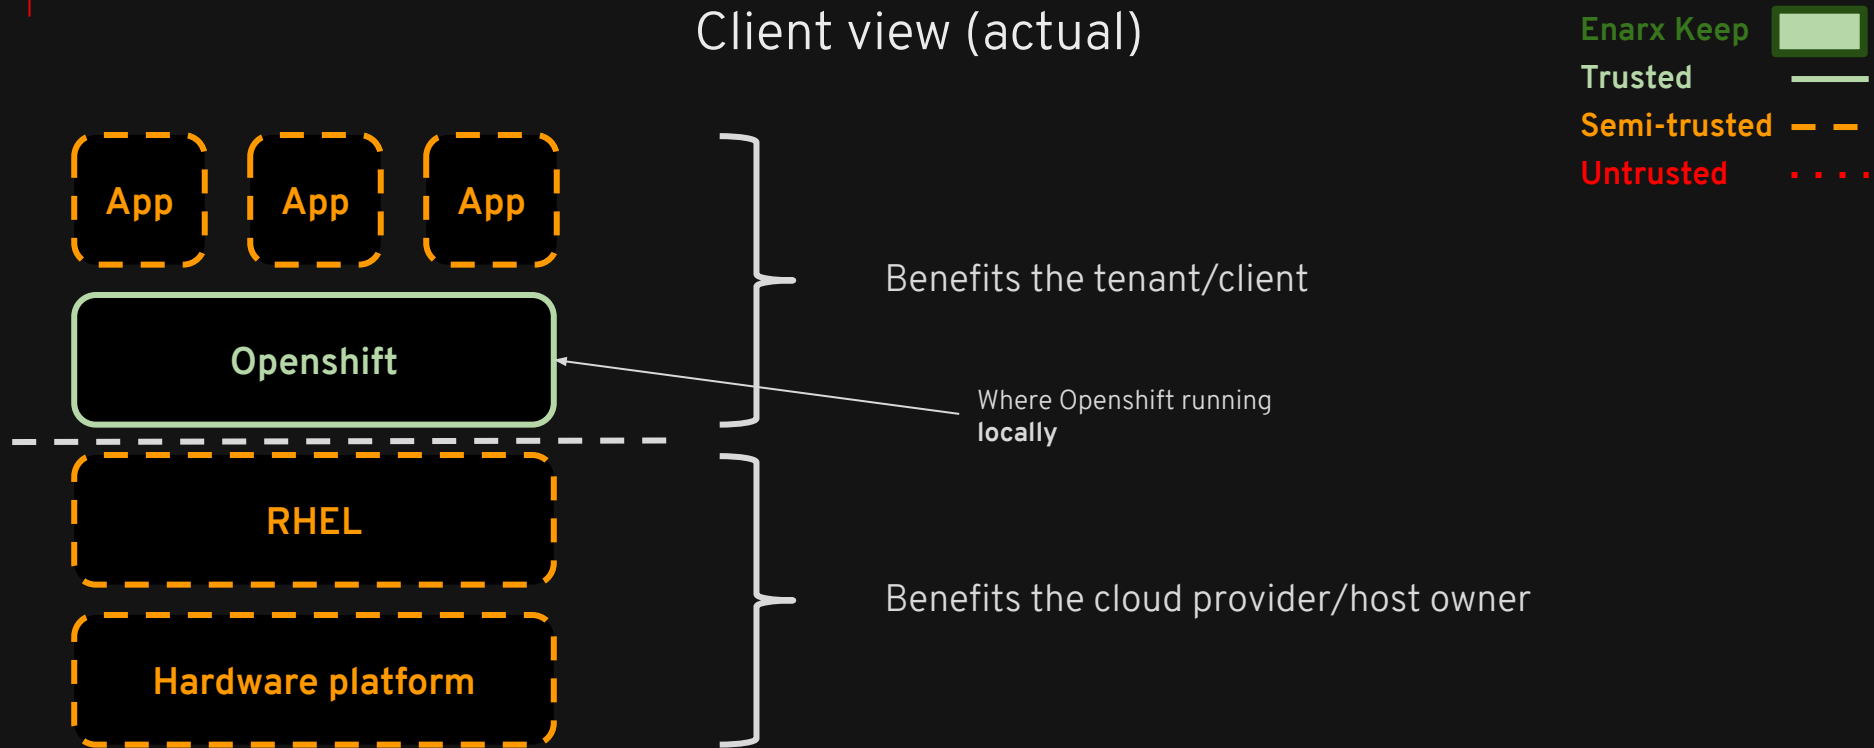
How trusted stacks help

Client view (actual)




How trusted stacks help

Client view (actual)




How trusted stacks help

Client view (with Enarx)




How trusted stacks help

Client view (with Enarx)



On which technology do I
build my application?

Introducing Enarx




Enarx Principles

1. We don't trust the host owner
2. We don't trust the host software
3. We don't trust the host users
4. We don't trust the host hardware
 - a. ... but we'll make an exception for CPU + firmware

Enarx Design Principles


1. Minimal Trusted Computing Base
2. Minimum trust relationships
3. Deployment-time portability
4. Network stack outside TCB
5. Security at rest, in transit and in use
6. Auditability
7. Open source
8. Open standards
9. Memory safety
10. No backdoors

Enarx Architecture



Keep - process or VM-based

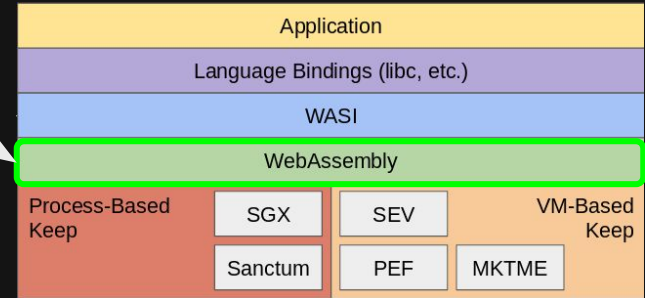
- Core Keep
- Platform-specific
 - Hardware (CPU): silicon vendor
 - Firmware: silicon vendor
 - Software: Enarx



Architecture varies between VM/Process-based platforms


WebAssembly (WASM)

- W3C standard
- Stack Machine ISA
- Sandboxed
- Supported by all browsers
- Exploding in the “serverless” space
- Implementations improving rapidly
 - cranelift and wasmtime



WebAssembly System API ([WASI](#))

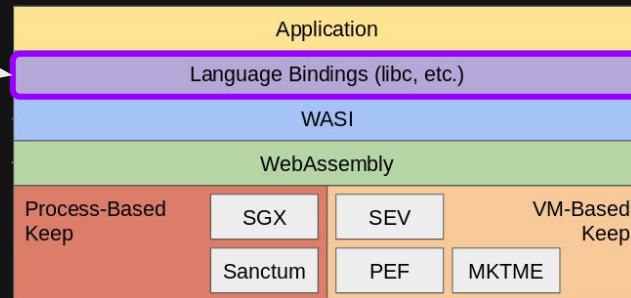
- W3C Standards Track
- Heavily inspired by a subset of POSIX
- Primary goals:
 - Portability
 - Security
- libc implementation on top
- Capability-based security:
 - No absolute resources
 - Think: `openat()` but not `open()`



Language Bindings (libc, etc.)


Compilation targets and includes, e.g.

- Rust: `--target wasm32-wasi`




Application

- Written by
 - Tenant (own development)
 - OR
 - 3rd party vendor
- Standard development tools
- Compiled to WebAssembly
- Using WASI interface




Enarx is a ~~Development~~ Deployment Framework



Enarx is a ~~Development~~ Deployment Framework

(Example components)



Best Practices? On By Default.


1. No Plaintext Networking (see **cipherpipe**)
2. No Plaintext Persisted Data
3. Independent Keep RNG
4. All Host APIs reviewed for data leakage

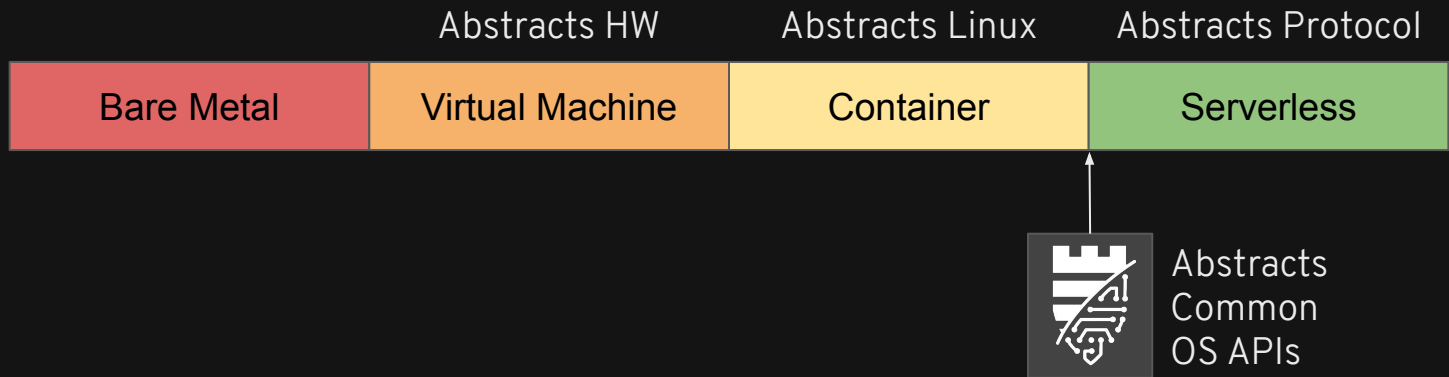
Bare Metal

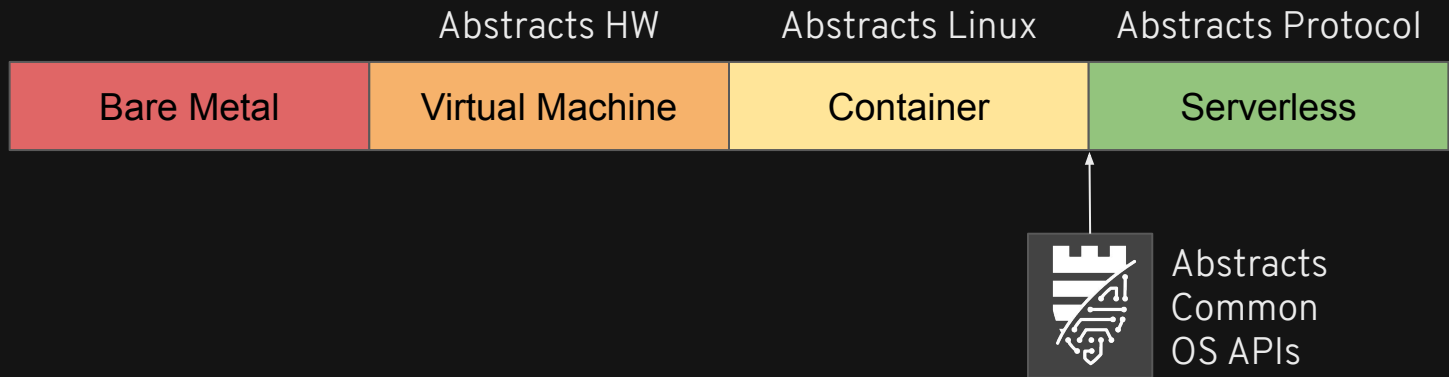
Virtual Machine

Container

Serverless








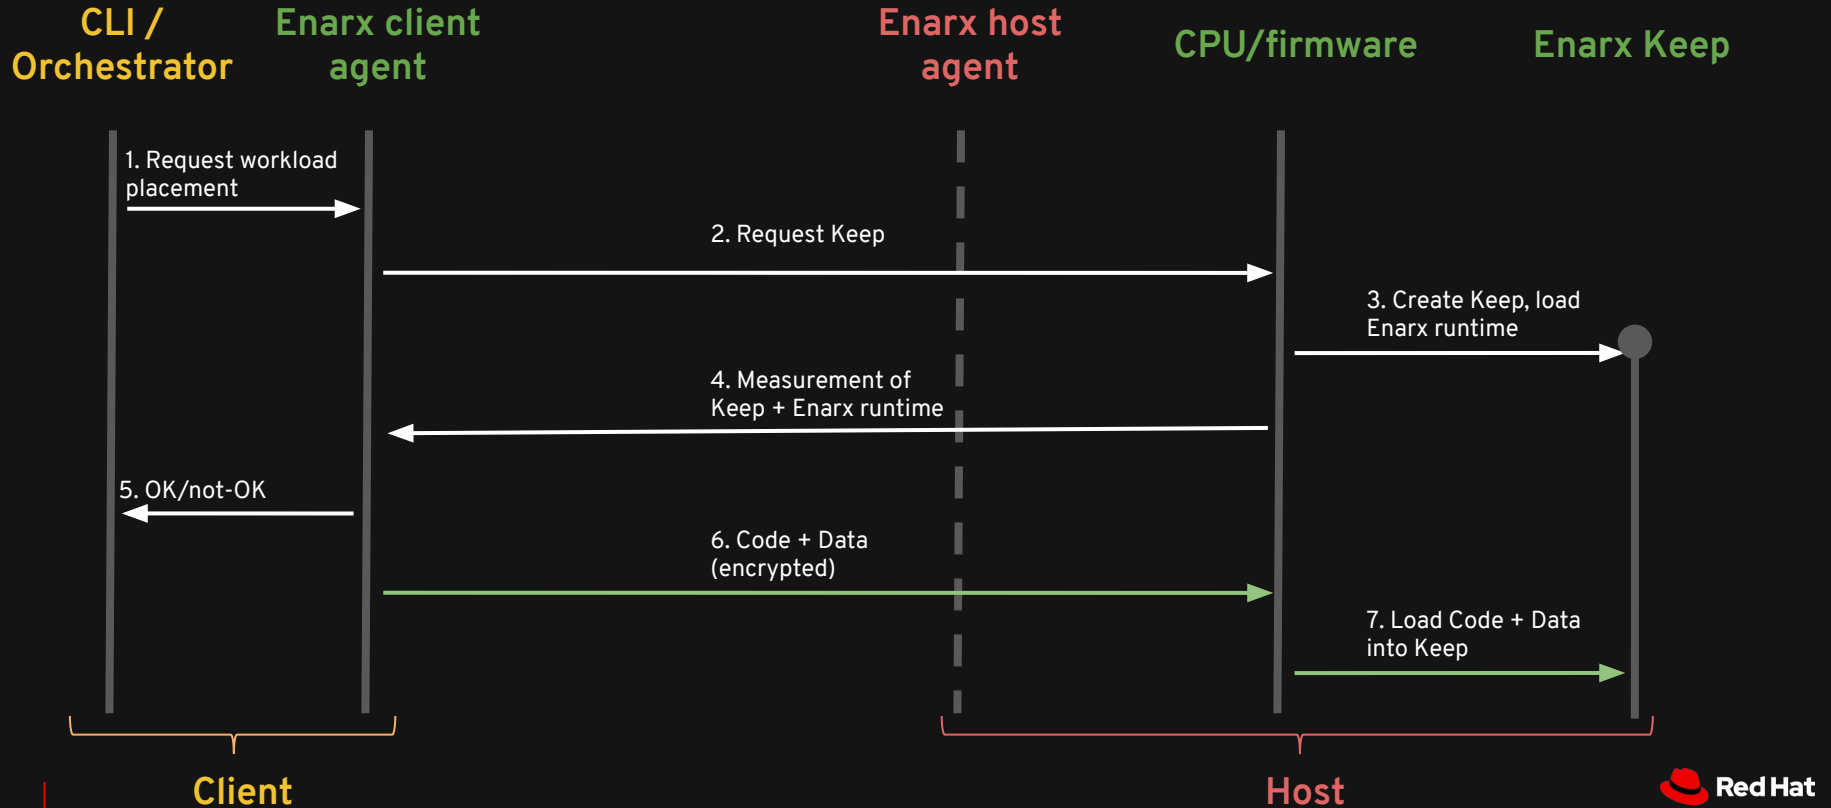
Just enough legacy support to enable trivial application portability.
Homogeneity to enable radical deployment-time portability.
No interfaces which accidentally leak data to the host.
Bridges process-based and VM-based TEE models.
No operating system to manage.

Process flow

Enarx architectural components




Enarx attestation process diagram




Enarx architectural components

Host

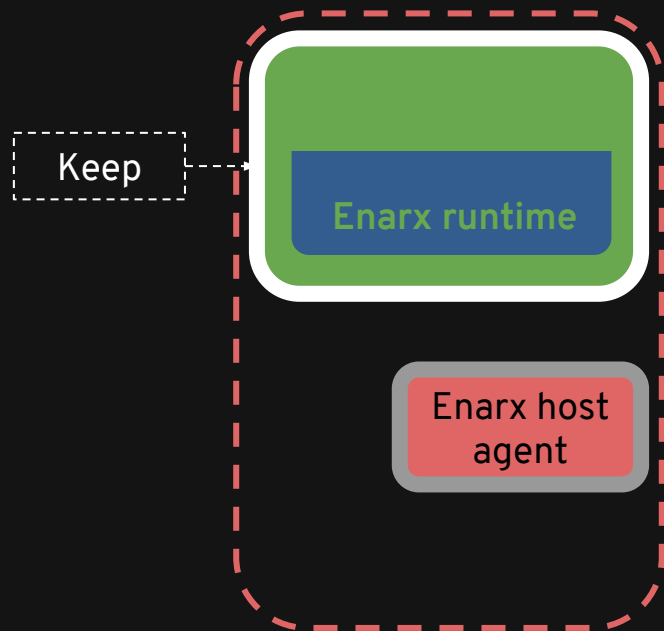


Client

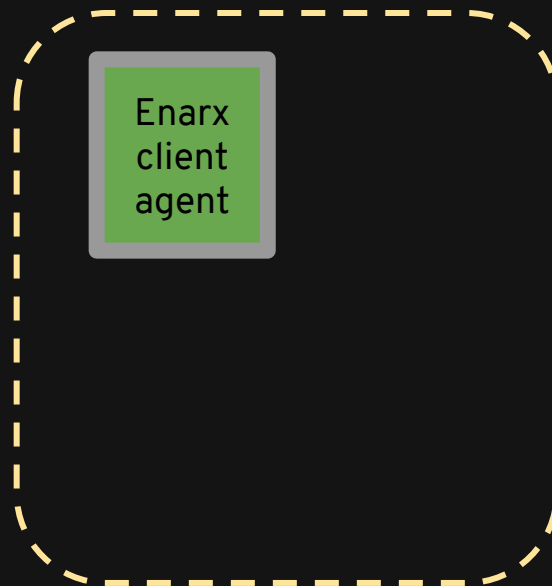


Enarx architectural components

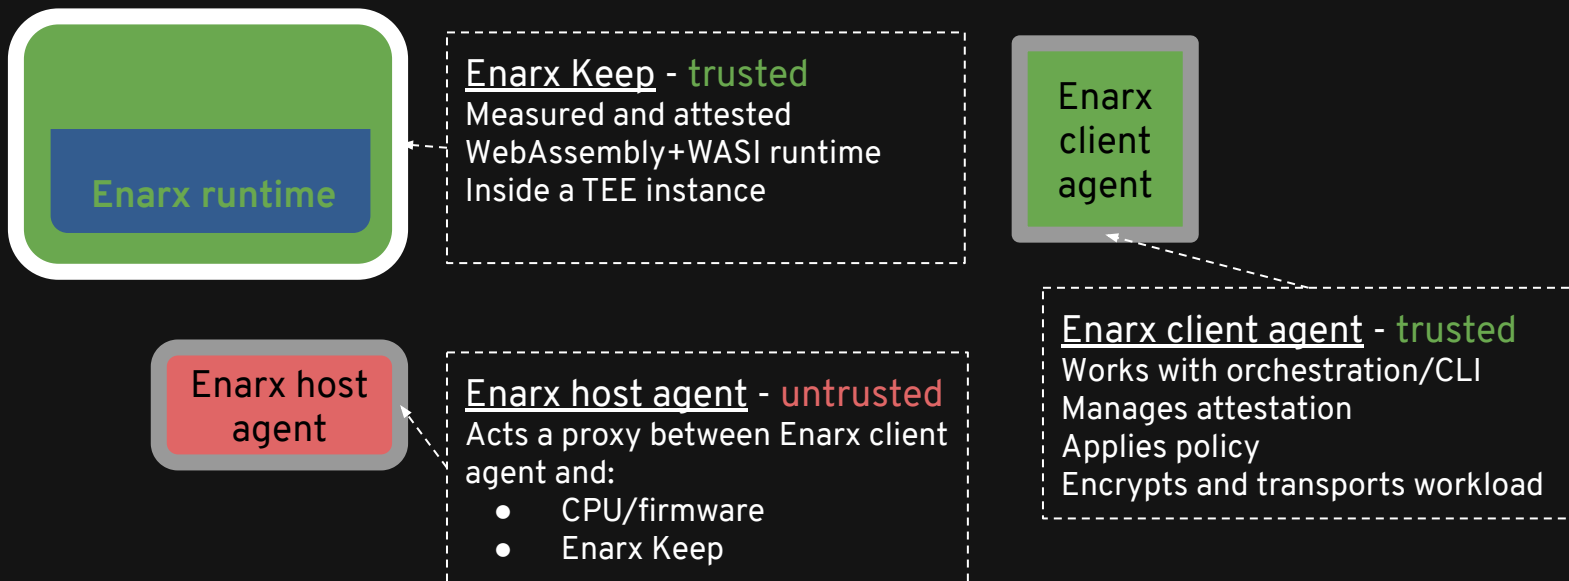
Host



Client





Enarx architectural components




Demo Time!

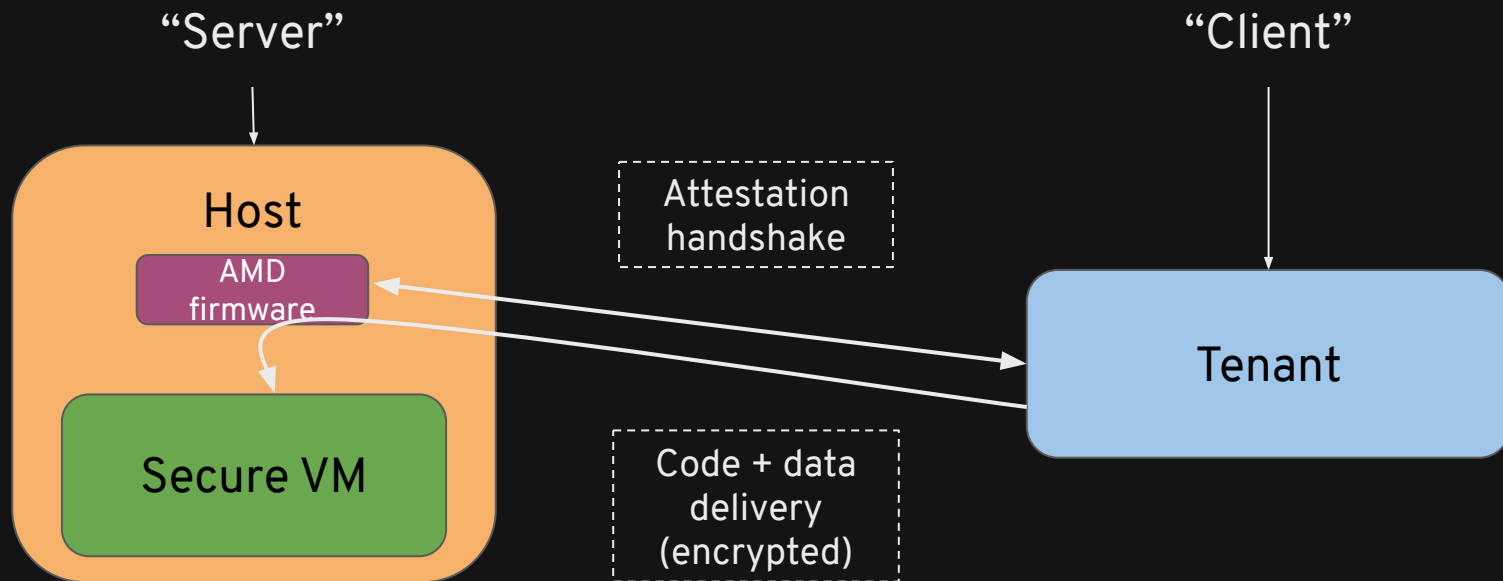
What will I see?




What will I see?

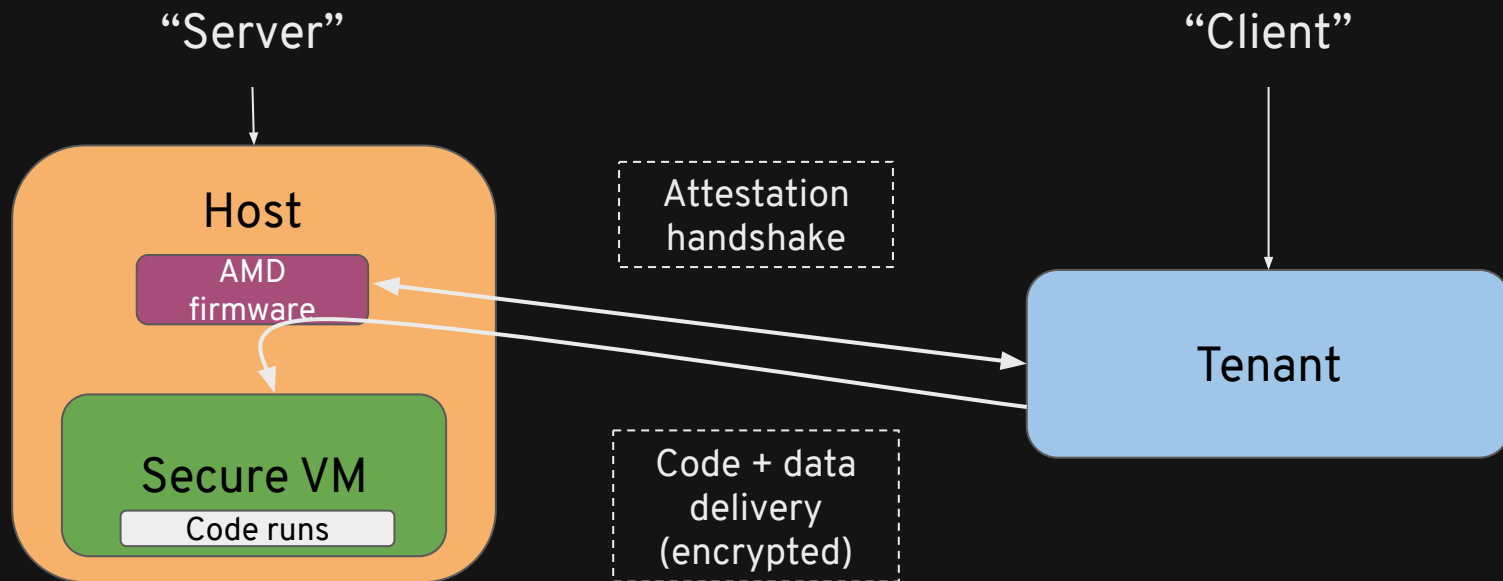


What will I see?






What did I see?

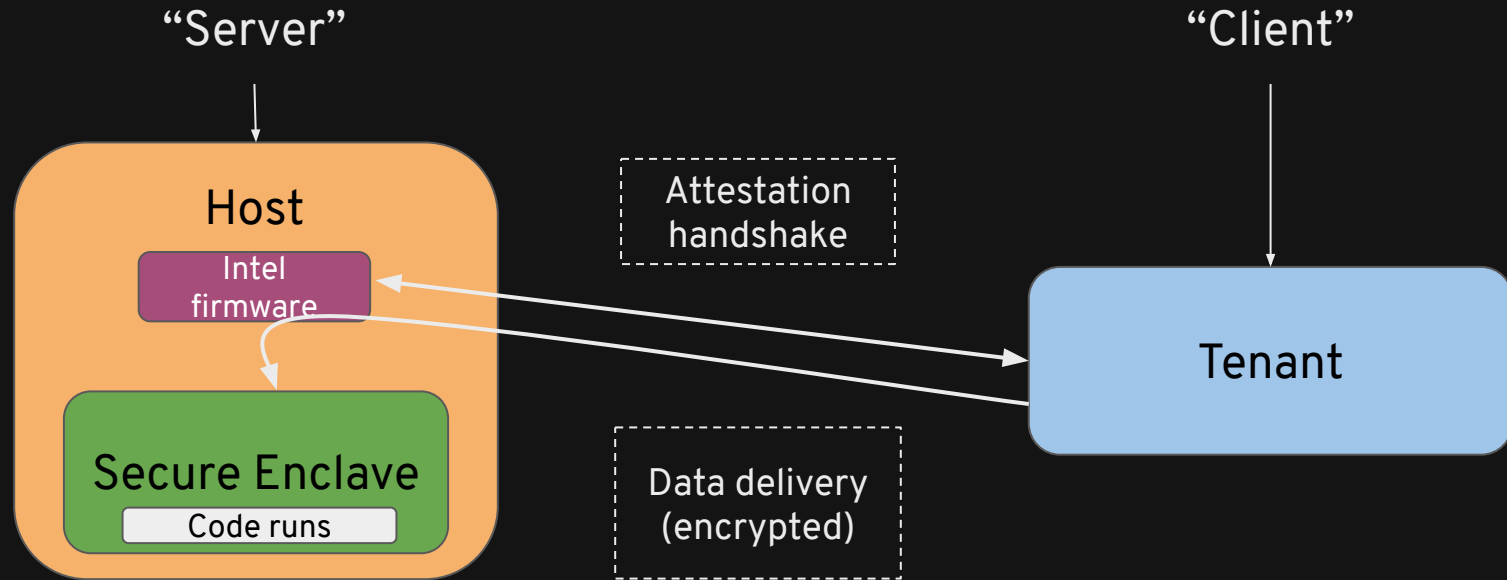


What did I see? (SEV)




```
enarxnuc@enarxnuc:~$ cd Repos/enarx/demo/intel-sgx/  
enarxnuc@enarxnuc:~/Repos/enarx/demo/intel-sgx$ time ./target/debug/attestation-tenant pck_cert_chain.pem 3 4  
CLIENT < SERVER: Quote (Attestation)  
CLIENT:      PCK cert chain OK  
CLIENT:      Quote signature OK  
CLIENT:      Attestation Key signature OK  
CLIENT:      Enclave report hash OK  
  
CLIENT:      Attestation Complete  
CLIENT > SERVER: Tenant PubKey and Encrypted Data  
  
7  
  
real    0m0.647s  
user    0m0.008s  
sys     0m0.004s  
enarxnuc@enarxnuc:~/Repos/enarx/demo/intel-sgx$
```

What did I see (SGX)?




What did I see?

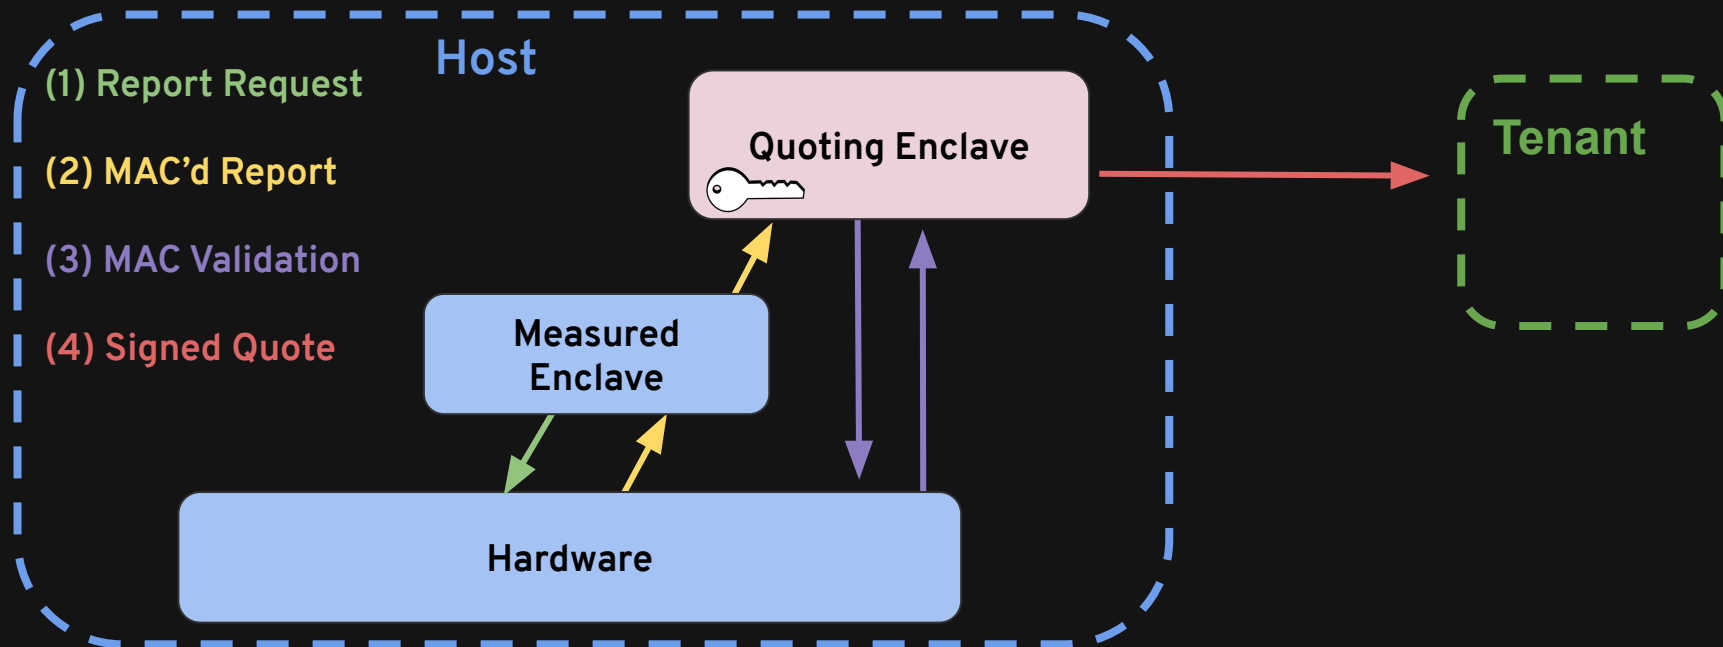


Intel SGX

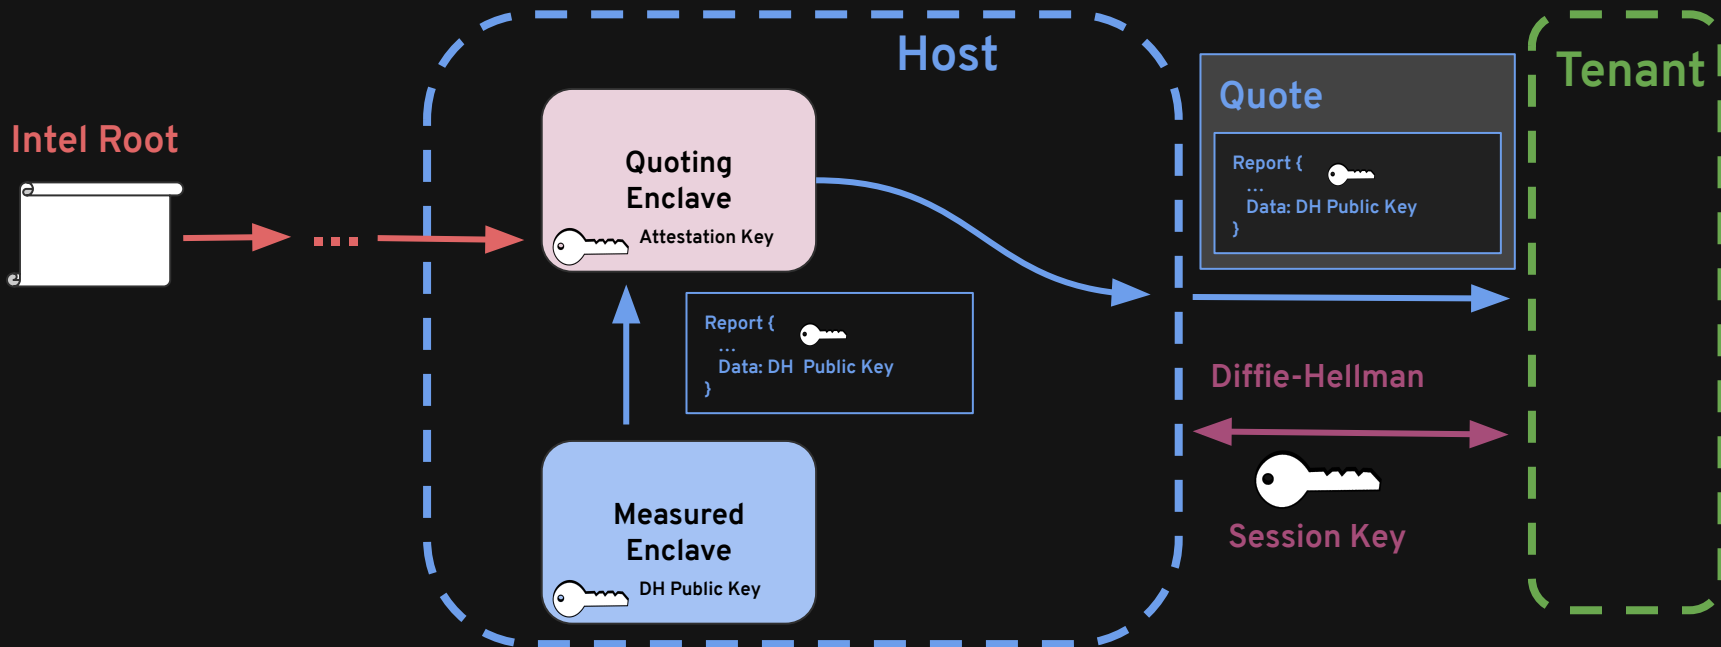
SGX Local Attestation



SGX Remote Attestation




SGX Secure Session

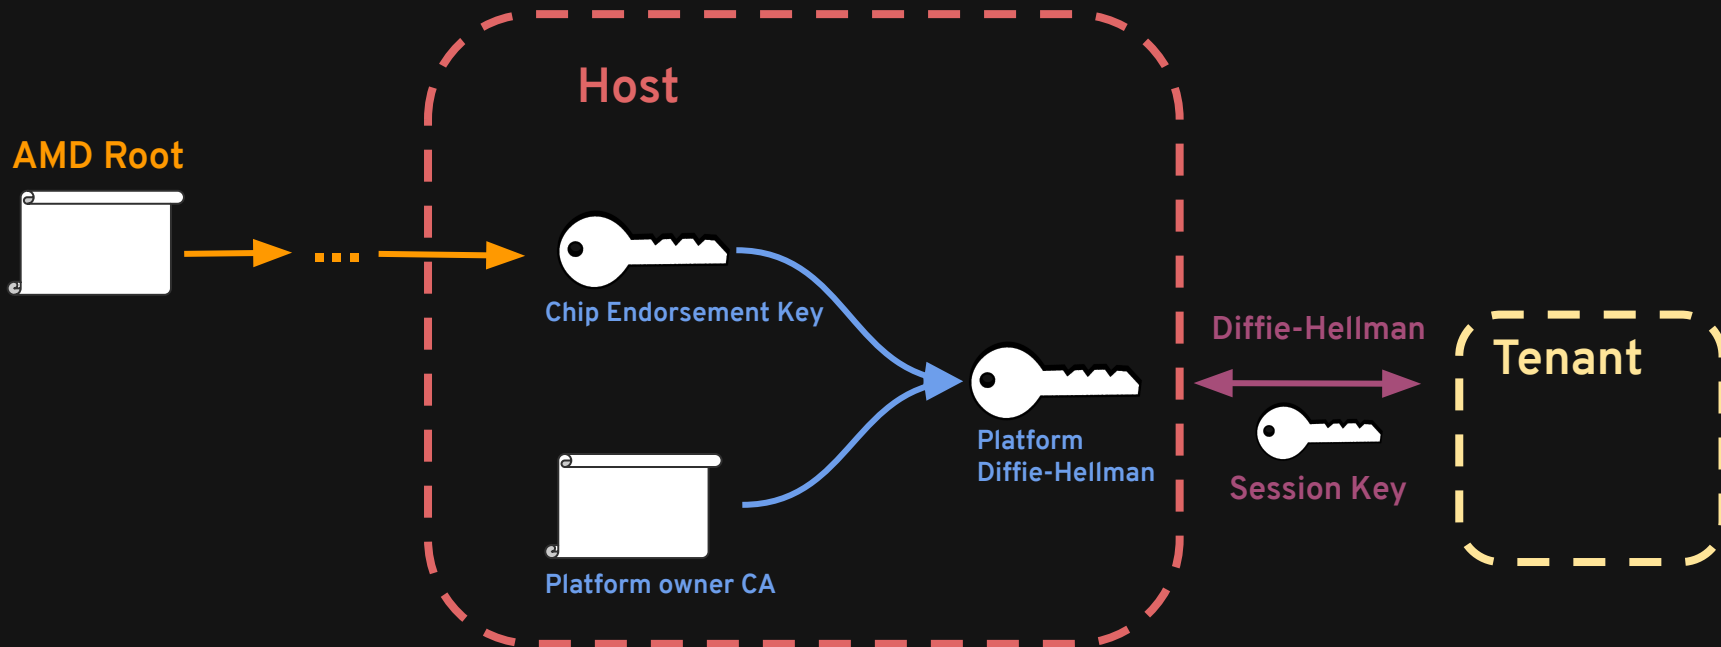


AMD SEV

SEV Attestation



SEV Secure Session

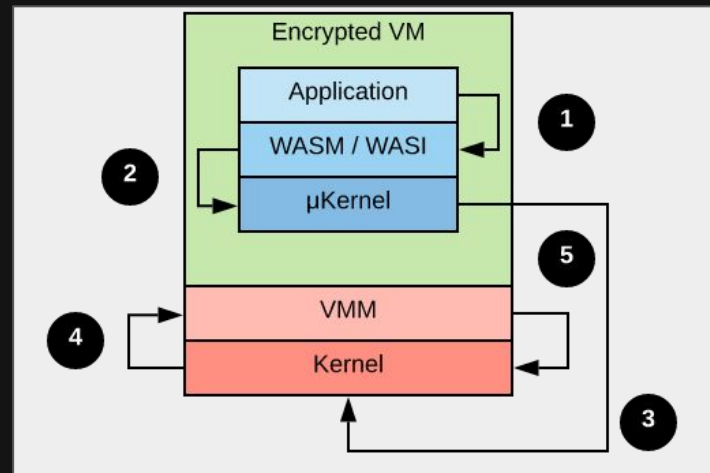


Enarx Virtualization Architecture

VM-based Keep


Syscall Propagation - VM-based

1. An Enarx application, compiled to WebAssembly, makes a WASI call, causing a transition from the JIT-compiled code into our guest userspace Rust code.




Syscall Propagation - VM-based

1. An Enarx application, compiled to WebAssembly, makes a WASI call, causing a transition from the JIT-compiled code into our guest userspace Rust code.
2. The hand-crafted Rust code translates the WASI call into a Linux `read()` syscall, leaving Ring 3 to jump into the μ Kernel, which handles some syscalls internally.




Syscall Propagation - VM-based

1. An Enarx application, compiled to WebAssembly, makes a WASI call, causing a transition from the JIT-compiled code into our guest userspace Rust code.
2. The hand-crafted Rust code translates the WASI call into a Linux `read()` syscall, leaving Ring 3 to jump into the μ Kernel, which handles some syscalls internally.
3. (Future work) Guest μ Kernel passes the syscall request to the host (Linux) kernel. As an optimization, some syscalls may be handled by the host (Linux) kernel directly.




Syscall Propagation - VM-based

1. An Enarx application, compiled to WebAssembly, makes a WASI call, causing a transition from the JIT-compiled code into our guest userspace Rust code.
2. The hand-crafted Rust code translates the WASI call into a Linux `read()` syscall, leaving Ring 3 to jump into the μ Kernel, which handles some syscalls internally.
3. (Future work) Guest μ Kernel passes the syscall request to the host (Linux) kernel. As an optimization, some syscalls may be handled by the host (Linux) kernel directly.
4. All syscalls which cannot be handled internally by the host kernel must cause a `vmexit` in the host VMM. Any syscalls which can be handled directly in the VMM are be handled immediately to avoid future context switches.




Syscall Propagation - VM-based

1. An Enarx application, compiled to WebAssembly, makes a WASI call, causing a transition from the JIT-compiled code into our guest userspace Rust code.
2. The hand-crafted Rust code translates the WASI call into a Linux `read()` syscall, leaving Ring 3 to jump into the μ Kernel, which handles some syscalls internally.
3. (Future work) Guest μ Kernel passes the syscall request to the host (Linux) kernel. As an optimization, some syscalls may be handled by the host (Linux) kernel directly.
4. All syscalls which cannot be handled internally by the host kernel must cause a `vmexit` in the host VMM. Any syscalls which can be handled directly in the VMM are be handled immediately to avoid future context switches.
5. In some cases, the VMM will have to re-enter the host kernel in order to fulfil the request. This is the slowest performance path and should be avoided wherever possible.



Syscall Propagation - VM-based

1. An Enarx application, compiled to WebAssembly, makes a WASI call, causing a transition from the JIT-compiled code into our guest userspace Rust code.
2. The hand-crafted Rust code translates the WASI call into a Linux `read()` syscall, leaving Ring 3 to jump into the μ Kernel, which handles some syscalls internally.
3. (Future work) Guest μ Kernel passes the syscall request to the host (Linux) kernel. As an optimization, some syscalls may be handled by the host (Linux) kernel directly.
4. All syscalls which cannot be handled internally by the host kernel must cause a `vmexit` in the host VMM. Any syscalls which can be handled directly in the VMM are be handled immediately to avoid future context switches.
5. In some cases, the VMM will have to re-enter the host kernel in order to fulfil the request. This is the slowest performance path and should be avoided wherever possible.



Enarx Status

Current Status

1. SEV: Fully attested demo w/ custom assembly.
 - a. Ketuvim: KVM library with SEV support
2. SGX: Fully attested demo w/ data delivery.
3. PEF: Ongoing discussions with POWER team.
4. WASM/WASI: Demo with some basic WASI functions.

Still To Do

<https://github.com/enarx/enarx/issues/1>

- Merge Ketuvim with rust-vmm
- Build:
 - Hypervisor
 - μ Kernel
 - WASI syscall propagation?
- Complete WASM JIT
- TLS networking stack
- Secure Clock?
- Research new platforms
- Openshift integration
- Much more...

Enarx Design Principles

1. Minimal Trusted Computing Base
2. Minimum trust relationships
3. Deployment-time portability
4. Network stack outside TCB
5. Security at rest, in transit and in use
6. Auditability
7. Open source
8. Open standards
9. Memory safety
10. No backdoors

We Need Your Help!

Website: <https://enarx.io>

Code: <https://github.com/enarx>

Gitter: <https://gitter.im/enarx/>

Master plan: <https://github.com/enarx/enarx/issues/1>

License: Apache 2.0

Language: Rust

Daily stand-ups open to all!
Check the website wiki for
details.

Questions?



<https://enarx.io>